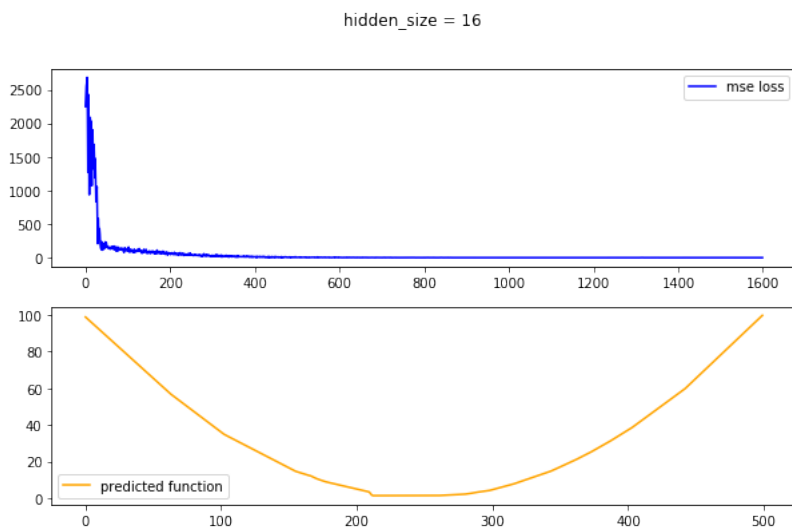
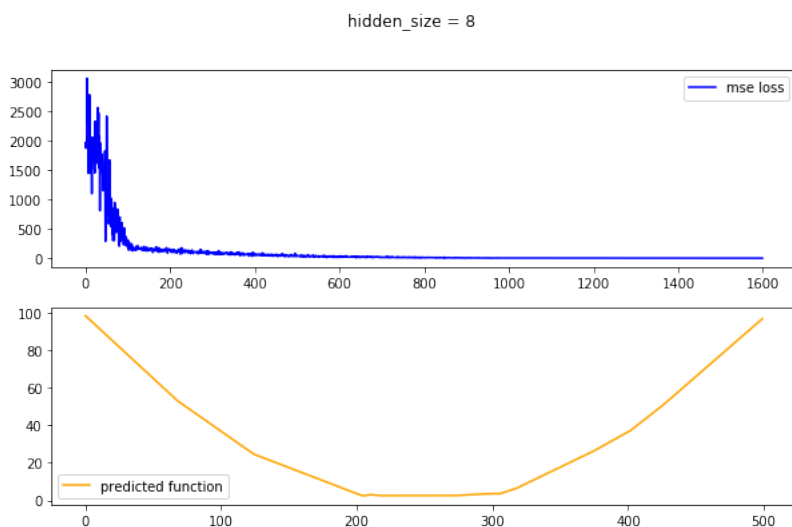



## Exercise 8

**Question 1.** How accurate is your learned model? How accurate it is within the range of  $[10,10]$ ? How about outside of the range? Can you notice any difference between models having different layer widths? What do you think that may have caused the difference?

*All code concerning this answer can be found in `question1.ipynb`.* My learned model is very accurate, achieving a minimum MSE loss of 0.0020633218. It attains this level of accuracy both inside and outside of its trained range of  $(-10, 10)$ . I noticed that as the number of hidden nodes increased, the prediction line better fit the line of the actual function. I assume this is because increasing the number of hidden nodes increases the expressive capability of the model to better approximate the function. I would assume this also results in the model having a higher variance.

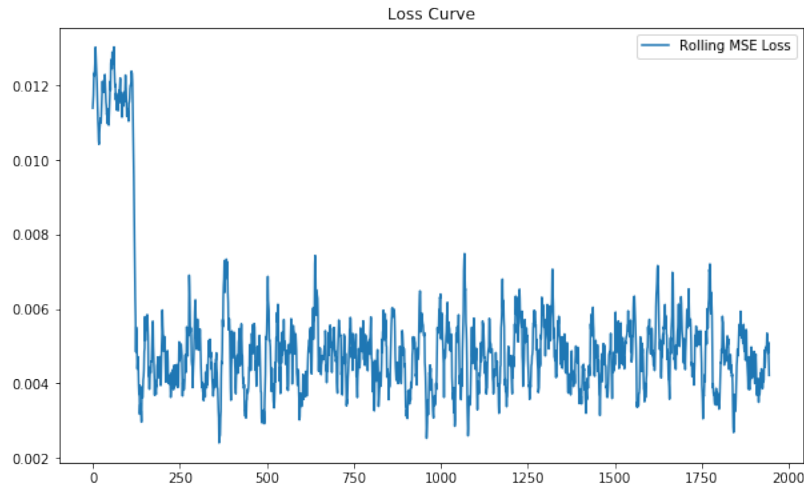




**Question 2.** Plot your learning curve (averaged over 10 trials) with confidence bands. How does it compare to your tabular methods from previous exercises?

*Please excuse that I could not run multiple trials (the plot below only includes one trial) as my machine could not train in a viable amount of time given my network architecture. All relevant code for this question can be found in `question2.ipynb`. The below plot shows the HuberLoss for my model over 2000 episodes. I used a learned feature method to represent the state to the DQN network, whereby I passed the board state (a 2D-matrix where walls are 2, current location is 1, and empty positions are 0) into two sequential*

convolutional layers. The performance of my DQN model was poor and performed worse than the tabular methods trialed in previous assignments. I have not yet pinpointed what is causing the model to not learn, but I assume more time altering my hyperparameters could possibly solve the issue.



### Question 3. Evaluate and tune your DQN on more environments

All relevant code for this question can be found in `question3a.py`. Helper functions and classes can be found in `cartpoledqn.py`, `replaymemory.py` and `cartpole-helpers.py`. Please activate the virtual environment requirements stored in `requirements.txt`. Below shows the live loss during training as demonstrated in the [pytorch documentation](#). The network parameters were three 2D convolutional layers followed by a single linear layer while the hyperparameters I use were batchsize = 128, gamma = 0.999, target-update = 10, and episodes = 500. I also gradually decayed epsilon to 0.05.

